

## Q1

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.graphics import tsaplots
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.holtwinters import SimpleExpSmoothing
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.exponential_smoothing.ets import ETSModel
from prophet import Prophet
import warnings
warnings.filterwarnings('ignore')
```

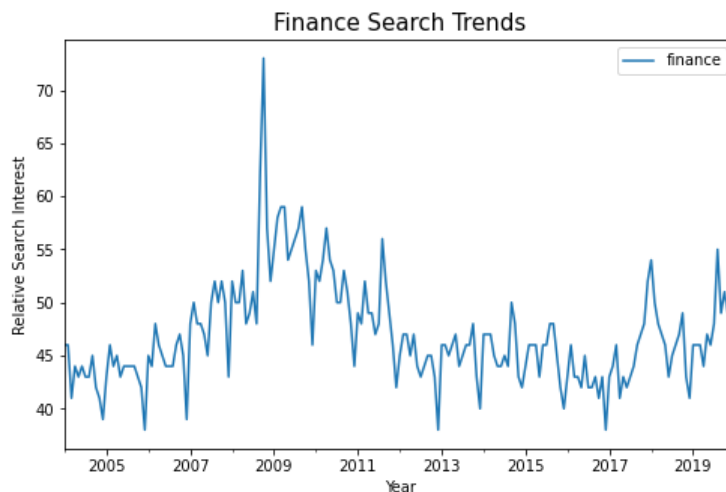
Importing plotly failed. Interactive plots will not work.

## A

```
In [2]: finance = pd.read_csv("finance_popularity.csv")
```

```
In [3]: def preprocess_and_plot(df, col = "month", freq='MS', title = "Finance Search Trends",
                                xlabel = 'Year', ylabel = 'Relative Search Interest', start = "", end = ""):
    df[col] = pd.date_range(start = df[col][0] if start == "" else start,
                            end = df[col].to_numpy()[-1] if end == "" else end, freq=freq)
    df.set_index(col, inplace = True)
    df.plot(figsize = (8,5))
    plt.title(title, fontsize=15)
    plt.xlabel(xlabel, fontsize=10)
    plt.ylabel(ylabel, fontsize=10)
    plt.show()
    return df
```

```
In [4]: finance = preprocess_and_plot(finance)
```




there seems to be a seasonality of 12 months according to the local bottom then a local peak at the start of each year, thus the data is not stationary

**B**

```
In [5]: def plot_diff_series(df = finance, periods=12, xlabel= "Year", ylabel = 'Relative Search Interest'):
df_diff=df.diff(periods=periods)
df_diff=df_diff.dropna()


df_diff.plot(figsize = (8,5))
plt.title("Differentiated Series",fontsize=15)
plt.xlabel(xlabel, fontsize=10)
plt.ylabel(ylabel, fontsize=10)
return df_diff
finance_year_diff = plot_diff_series()
```




we removed the yearly seasonality and also we can see that there is no trend as the data is distributed around 0 therefore the data is stationary

**C**

```
In [6]: tsaplots.plot_acf(finance_year_diff);
```



```
In [7]: tsaplots.plot_pacf(finance_year_diff);
```



the pacf is around zero at 3 or more lags, thus we can assume that a good model would be  $AR(2)$  or  $SARIMA(2, 0, 0)(1, 1, 0)_{12}$

D

$SARIMA(2, 0, 0)(1, 1, 0)_{12}$

```
In [8]: def fit_summarize_sarima(data ,order, seasonal_order = (0,0,0,0)):
        sarima_model = SARIMAX(data, order=order,seasonal_order=seasonal_order, freq="MS")
        sarima_model_fit = sarima_model.fit(dis= False)
        display(sarima_model_fit.summary())
        fit_summarize_sarima(finance, (2, 0, 0), (1,1,0,12))
```

SARIMAX Results

Dep. Variable:	finance				No. Observations:	193
Model:	SARIMAX(2, 0, 0)x(1, 1, 0, 12)				Log Likelihood	-464.106
Date:	Wed, 14 Feb 2024				AIC	936.212
Time:	13:03:29				BIC	949.006
Sample:	01-01-2004				HQIC	941.399
	- 01-01-2020					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5935	0.044	13.570	0.000	0.508	0.679
ar.L2	0.1923	0.063	3.075	0.002	0.070	0.315
ar.S.L12	-0.4594	0.044	-10.479	0.000	-0.545	-0.373
sigma2	9.6836	0.661	14.648	0.000	8.388	10.979
Ljung-Box (L1) (Q):	0.14	Jarque-Bera (JB):	175.39			
Prob(Q):	0.71	Prob(JB):	0.00			
Heteroskedasticity (H):	0.53	Skew:	0.76			
Prob(H) (two-sided):	0.01	Kurtosis:	7.57			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

AR(2)

```
In [9]: fit_summarize_sarima(finance, (2, 0, 0))
```

SARIMAX Results

Dep. Variable:	finance	No. Observations:	193
Model:	SARIMAX(2, 0, 0)	Log Likelihood	-514.193
Date:	Wed, 14 Feb 2024	AIC	1034.386
Time:	13:03:29	BIC	1044.174
Sample:	01-01-2004	HQIC	1038.350
	- 01-01-2020		
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.8166	0.038	21.481	0.000	0.742	0.891
ar.L2	0.1813	0.039	4.674	0.000	0.105	0.257
sigma2	11.7378	0.776	15.125	0.000	10.217	13.259

Ljung-Box (L1) (Q):	0.56	Jarque-Bera (JB):	66.56
Prob(Q):	0.45	Prob(JB):	0.00
Heteroskedasticity (H):	0.43	Skew:	0.35
Prob(H) (two-sided):	0.00	Kurtosis:	5.79

Warnings:  
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

E

the  $SARIMA(2, 0, 0)(1, 1, 0)_{12}$  model has a better AIC and BIC so we will choose this as our model, lets try to forecast the last year data (without fitting on it) to check how the model works.

```
In [10]: def forecast_with_sarimax(time_series, forecast_periods, order, seasonal_order, xlabel = 'Year',
                                     ylabel = 'Relative Search Interest', freq='MS', plot=True, offset = 0, title =
                                     "", forecast_next = False):
    if forecast_next:
        sarima_model = SARIMAX(time_series, order=order, seasonal_order=seasonal_order)
    else:
        sarima_model = SARIMAX(time_series[:-forecast_periods - offset], order=order, seasonal_order=seasonal_order)
    sarima_model_fit = sarima_model.fit(dis=False)

    next_periods_forecast = sarima_model_fit.get_forecast(steps=forecast_periods)
    next_periods_ci = next_periods_forecast.conf_int()

    if forecast_next:
        future_dates = pd.date_range(start=time_series.index[-1], periods=forecast_periods, freq=freq)
    else:
        future_dates = pd.date_range(start=time_series.index[-forecast_periods - offset], periods=forecast_periods, freq=freq)
    next_periods_ci.index = future_dates


    futureDF_forecast = pd.DataFrame(next_periods_forecast.predicted_mean, index=future_dates)

    if plot:
        if offset == 0:
            ax = time_series.plot(label='Observed', figsize=(10, 5))
        else:
            ax = time_series[:-offset].plot(label='Observed', figsize=(10, 5))
        futureDF_forecast.plot(ax=ax, label='Forecast', legend=True)

        ax.fill_between(next_periods_ci.index,
                        next_periods_ci.iloc[:, 0],
                        next_periods_ci.iloc[:, 1], color='k', alpha=.1)

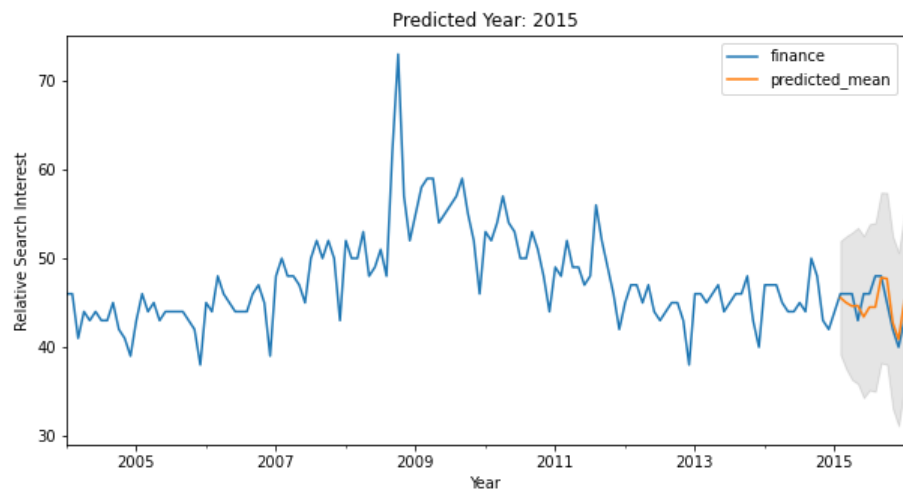

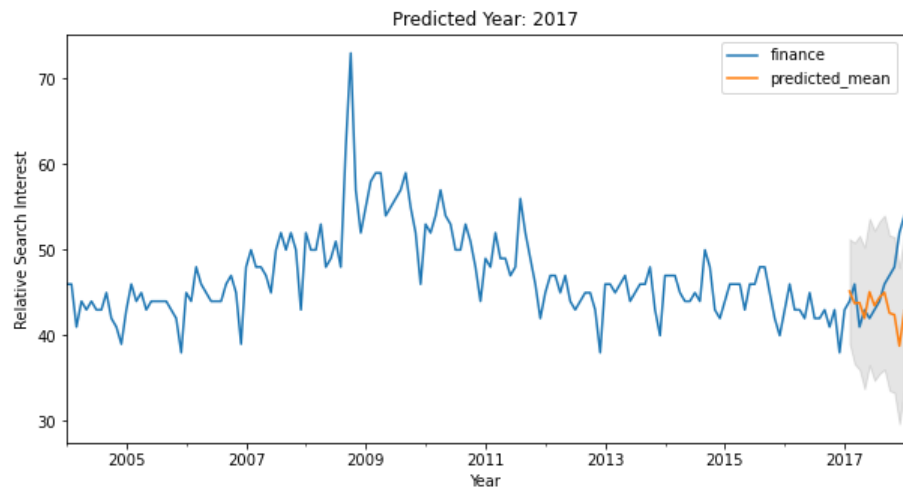

        ax.set_title(title)
        ax.set_xlabel(xlabel)
        ax.set_ylabel(ylabel)
        plt.legend()
        plt.show()
    return futureDF_forecast, next_periods_ci
```


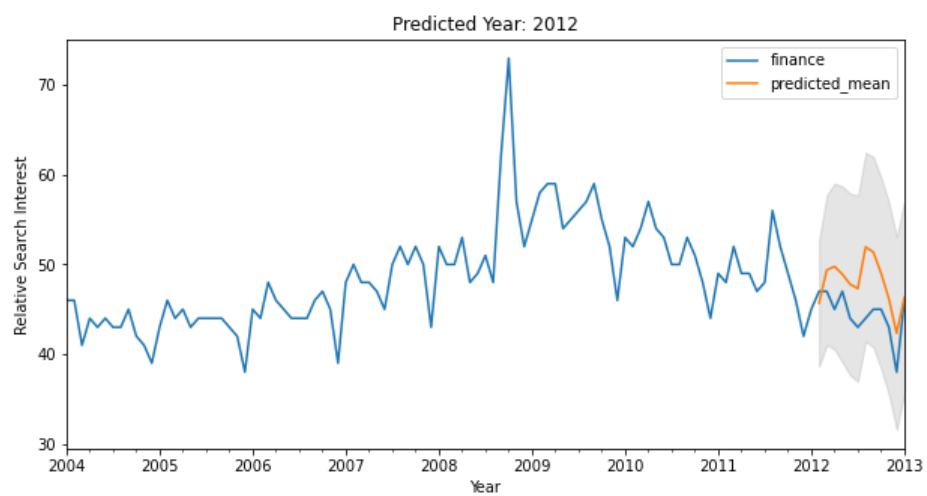


```
future_finance_forecast, next_periods_ci = forecast_with_sarimax(finance, 12, order=(2, 0, 0), seasonal_order=(1, 1, 0, 12))
```



a one data sample for checking the model is rather low, lets create a loop to try several years

```
In [11]: FORECAST = 12
OFFSET = 12
for year, OFFSET in zip(range(2018, 2010, -1), range(12, 9*12,12)):
    future_finance_forecast, next_periods_ci = forecast_with_sarimax(finance, FORECAST, order=(2, 0, 0),seasonal_order=(1, 1, 0,12),
                                                                    offset = OFFSET, title = f"Predicted Year: {year}")
```







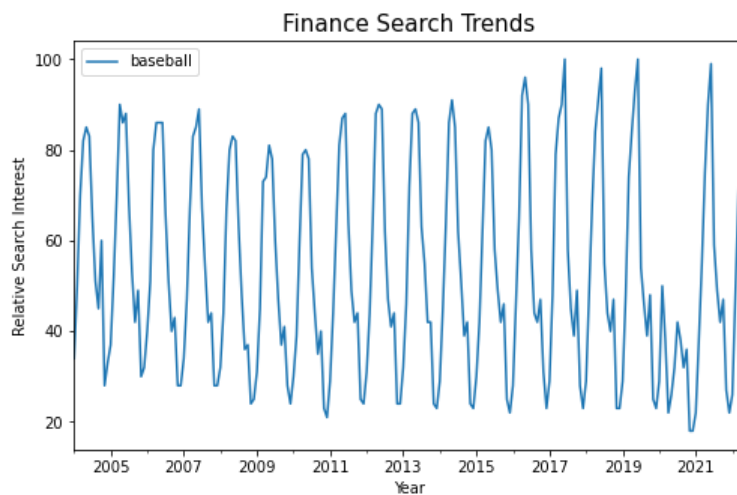
the model seem to work quite well, i tried some different models but did not get much better AIC, BIC and plots

## Q2

### A

```
In [12]: baseball = pd.read_csv("baseball_popularity.csv")
```

```
In [13]: baseball = preprocess_and_plot(baseball)
```

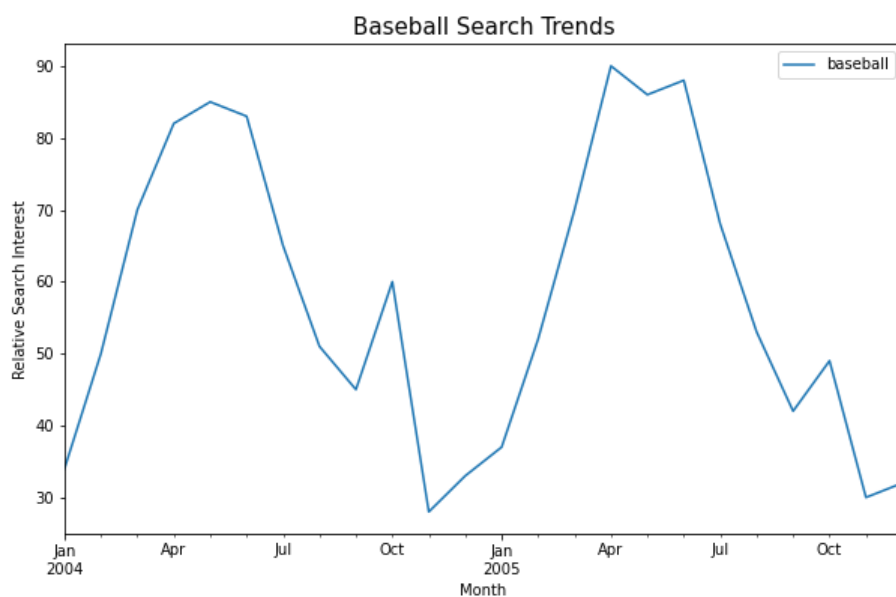


this time series also looks seasonal without a trend, thus it is not stationary. baseball season usually starts at march\april and ends at october, so we can expect a 12 month seasonality

taking a look at the first 2 years

```
In [14]: baseball[:24].plot(figsize = (10,6))
plt.title("Baseball Search Trends", fontsize=15)
plt.xlabel('Month', fontsize=10)
plt.ylabel('Relative Search Interest', fontsize=10)
```

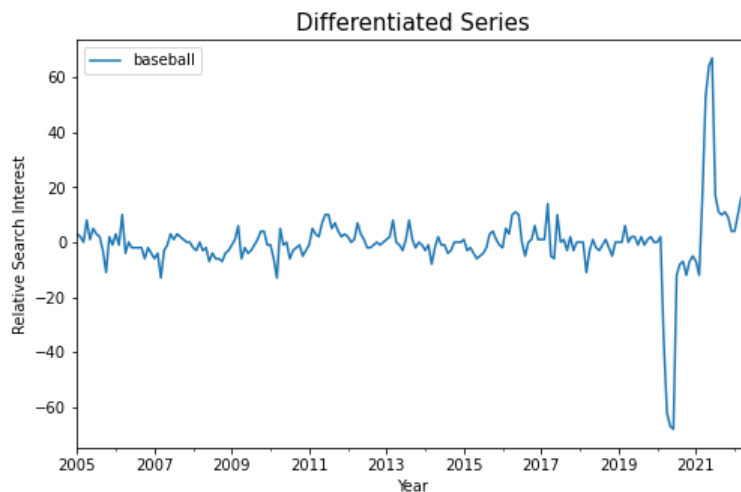
```
Out[14]: Text(0, 0.5, 'Relative Search Interest')
```



we can see a peak at the start of the season that a slowly decays, until another local peak at october probably due to the finals.

**B**


```
In [15]: baseball_year_diff = plot_diff_series(baseball)
```




looks stationary, except for the major noise at 2020-2021 (corona)?

**C**

```
In [16]: tsaplots.plot_acf(baseball_year_diff);
```



```
In [17]: tsaplots.plot_pacf(baseball_year_diff);
```



the graphs suggests a strong correlation between the last 2 lags and maybe the third lag but also for the lags a year ago, so we can assume that a good model would be  $AR(2)$  or  $AR(3)$   $SARIMA(2, 0, 0)(1, 1, 0)_{12}$  or  $SARIMA(2, 0, 0)(2, 1, 0)_{12}$ , or even maybe Holt Winter's Exponential Smoothing with 12 months seasonality

D

$SARIMA(2, 0, 0)(1, 1, 0)_{12}$

```
In [18]: fit_summarize_sarima(baseball, (2,0,0), (1,1,0,12))
```

SARIMAX Results

Dep. Variable:	baseball	No. Observations:	221			
Model:	SARIMAX(2, 0, 0)x(1, 1, 0, 12)	Log Likelihood	-696.006			
Date:	Wed, 14 Feb 2024	AIC	1400.013			
Time:	13:03:32	BIC	1413.382			
Sample:	01-01-2004	HQIC	1405.418			
	- 05-01-2022					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.9349	0.034	27.439	0.000	0.868	1.002
ar.L2	-0.2219	0.041	-5.461	0.000	-0.302	-0.142
ar.S.L12	-0.5296	0.035	-15.308	0.000	-0.597	-0.462
sigma2	44.6508	2.200	20.292	0.000	40.338	48.963
Ljung-Box (L1) (Q):	0.03	Jarque-Bera (JB):	759.63			
Prob(Q):	0.87	Prob(JB):	0.00			
Heteroskedasticity (H):	5.23	Skew:	-0.07			
Prob(H) (two-sided):	0.00	Kurtosis:	12.34			

Warnings:  
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

*SARIMA*(2, 0, 0)(2, 1, 0)<sub>12</sub>

```
In [19]: fit_summarize_sarima(baseball, (2,0,0), (2,1,0,12))
```

SARIMAX Results

Dep. Variable:	baseball	No. Observations:	221			
Model:	SARIMAX(2, 0, 0)x(2, 1, 0, 12)	Log Likelihood	-690.039			
Date:	Wed, 14 Feb 2024	AIC	1390.078			
Time:	13:03:33	BIC	1406.790			
Sample:	01-01-2004	HQIC	1396.835			
	- 05-01-2022					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.8810	0.036	24.285	0.000	0.810	0.952
ar.L2	-0.1692	0.045	-3.794	0.000	-0.257	-0.082
ar.S.L12	-0.6692	0.057	-11.748	0.000	-0.781	-0.558
ar.S.L24	-0.3245	0.067	-4.821	0.000	-0.456	-0.193
sigma2	41.7372	1.940	21.519	0.000	37.936	45.539
Ljung-Box (L1) (Q):	0.02	Jarque-Bera (JB):	1060.91			
Prob(Q):	0.89	Prob(JB):	0.00			
Heteroskedasticity (H):	5.87	Skew:	-0.60			
Prob(H) (two-sided):	0.00	Kurtosis:	13.97			

Warnings:  
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

*AR*(2)

```
In [20]: fit_summarize_sarima(baseball, (2,0,0))
```

SARIMAX Results

Dep. Variable:	baseball	No. Observations:	221			
Model:	SARIMAX(2, 0, 0)	Log Likelihood	-885.737			
Date:	Wed, 14 Feb 2024	AIC	1777.474			
Time:	13:03:33	BIC	1787.669			
Sample:	01-01-2004	HQIC	1781.591			
	- 05-01-2022					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	1.3758	0.075	18.277	0.000	1.228	1.523
ar.L2	-0.4219	0.083	-5.105	0.000	-0.584	-0.260
sigma2	174.7732	19.362	9.027	0.000	136.825	212.722
Ljung-Box (L1) (Q):	0.70	Jarque-Bera (JB):	62.03			
Prob(Q):	0.40	Prob(JB):	0.00			
Heteroskedasticity (H):	1.81	Skew:	-1.11			
Prob(H) (two-sided):	0.01	Kurtosis:	4.34			

Warnings:  
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

$AR(3)$ 

```
In [21]: fit_summarize_sarima(baseball, (3,0,0))
```

## SARIMAX Results

**Dep. Variable:** baseball **No. Observations:** 221  
**Model:** SARIMAX(3, 0, 0) **Log Likelihood** -885.468  
**Date:** Wed, 14 Feb 2024 **AIC** 1778.937  
**Time:** 13:03:33 **BIC** 1792.529  
**Sample:** 01-01-2004 **HQIC** 1784.425  
- 05-01-2022  
**Covariance Type:** opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	1.3550	0.081	16.675	0.000	1.196	1.514
ar.L2	-0.3535	0.148	-2.392	0.017	-0.643	-0.064
ar.L3	-0.0499	0.101	-0.496	0.620	-0.247	0.147
sigma2	174.2810	19.538	8.920	0.000	135.987	212.575

**Ljung-Box (L1) (Q):** 0.13 **Jarque-Bera (JB):** 62.61  
**Prob(Q):** 0.72 **Prob(JB):** 0.00  
**Heteroskedasticity (H):** 1.83 **Skew:** -1.11  
**Prob(H) (two-sided):** 0.01 **Kurtosis:** 4.37


## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Prophet


```
In [22]: baseball_prophet = baseball.reset_index().rename({"month": "ds", "baseball": "y"}, axis =1)
m = Prophet()
m.fit(baseball_prophet)
future_dates = m.make_future_dataframe(periods=60, freq='MS')
forecast = m.predict(future_dates)
forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']]
m.plot_components(forecast);
```

13:03:33 - cmdstanpy - INFO - Chain [1] start processing  
13:03:33 - cmdstanpy - INFO - Chain [1] done processing



we can see that there is indeed yearly seasonality April-June seem to be the time when people usually search baseball. we also see a linear decreasing trend

```
In [23]: m.plot(forecast);
```



## HWES

```

In [24]: import matplotlib.pyplot as plt

def plot_observed_predictions(data, preds, forecast_len, ylabel='', xlabel=''):
    fig, ax = plt.subplots(figsize=(10, 6))
    ax.plot(data.index, data, label="Observed")
    ax.plot(data.index, preds[:-forecast_len], color="orange", ls = "-.", label="Predictions")
    ax.set_xlabel(xlabel)
    ax.set_ylabel(ylabel)
    plt.legend()
    plt.show()

def plot_forecast_with_confidence_intervals(data, preds, upper_ci, lower_ci, ylabel='', xlabel=''):
    ax = data.plot(label="Original series", figsize=(10, 6))
    preds.plot(label="Forecast", color="orange", ax=ax)
    ax.fill_between(upper_ci.index, upper_ci, lower_ci, color='k', alpha=.2)
    ax.set_xlabel(xlabel)
    ax.set_ylabel(ylabel)
    plt.legend()
    plt.show()

def simulate_ets_predictions(data, column_name, seasonal='add', trend='add', seasonal_periods=12, n_repetitions=500,
                             forecast_len=60, plot_obs_pred=True, ylabel='', xlabel="", plot_org_forecast=True):
    ets_model = ETSModel(
        endog=data[column_name],
        seasonal=seasonal,
        trend=trend,
        seasonal_periods=seasonal_periods)
    ets_result = ets_model.fit()

    # Simulate predictions.
    n_steps_prediction = data[column_name].shape[0]

    df_simul = ets_result.simulate(
        nsimulations=n_steps_prediction+forecast_len,
        repetitions=n_repetitions,
        anchor='start',
    )
    preds = df_simul.mean(axis=1)
    # Calculate confidence intervals.
    upper_ci = df_simul.quantile(q=0.95, axis='columns')
    lower_ci = df_simul.quantile(q=0.05, axis='columns')



    if plot_obs_pred:
        plot_observed_predictions(data[column_name], preds, forecast_len, ylabel, xlabel)

    if plot_org_forecast:
        plot_forecast_with_confidence_intervals(data[column_name], preds[:-forecast_len:],
                                                upper_ci[:-forecast_len:], lower_ci[:-forecast_len:], ylabel, xlabel)

    return preds, lower_ci, upper_ci

preds, lower_ci, upper_ci = simulate_ets_predictions(baseball, "baseball",
                                                    ylabel = 'Baseball reletive search index', xlabel="year")

```




E

$SARIMA(2, 0, 0)(2, 1, 0)_{12}$  has the best AIC and BIC (excluding prophet and HWES)

prophet and HWES have a similar fit based on the graphs

```
In [25]: _, _ = forecast_with_sarimax(baseball, 60, order=(2, 0, 0), seasonal_order=(2, 1, 0, 12), ylabel = "baseball search index")
```






$SARIMA(2, 0, 0)(2, 1, 0)_{12}$  also seem like a good fit by the graph and a little better than the HWES and PROPHET so we will chose this as our model

## F


```
In [26]: future_finance_forecast, next_periods_ci = forecast_with_sarimax(baseball, 60, order=(2, 0, 0),seasonal_order=(2, 1, 0,12),
                                ylabel = "baseball search index", forecast
                                _next = True)
```



## Q3

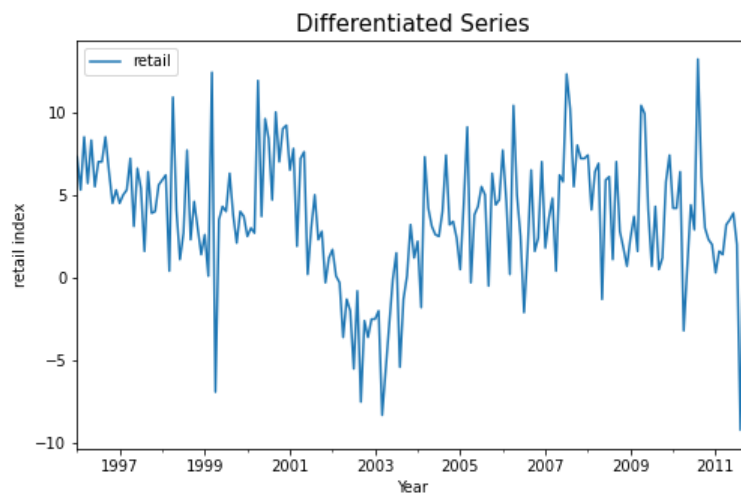
## A

```
In [27]: retail = pd.read_csv("retail.csv")
retail = preprocess_and_plot(retail, start = '1995-01-01', end = '2011-10-01',ylabel = "retail index")
```



we can see an uptrend with some seasonlity, lets fit hte sarima model and check how it fits the data

```
In [28]: _ = plot_diff_series(retail, ylabel = "retail index")
```



```
In [29]: fit_summarize_sarima(retail, (0,1,1), (0,1,1,12))
```

#### SARIMAX Results

<b>Dep. Variable:</b>	retail	<b>No. Observations:</b>	202
<b>Model:</b>	SARIMAX(0, 1, 1)x(0, 1, 1, 12)	<b>Log Likelihood</b>	-453.669
<b>Date:</b>	Wed, 14 Feb 2024	<b>AIC</b>	913.338
<b>Time:</b>	13:03:36	<b>BIC</b>	923.063
<b>Sample:</b>	01-01-1995	<b>HQIC</b>	917.278
	- 10-01-2011		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	-0.7060	0.052	-13.462	0.000	-0.809	-0.603
<b>ma.S.L12</b>	-0.7814	0.065	-11.967	0.000	-0.909	-0.653
<b>sigma2</b>	6.6805	0.569	11.732	0.000	5.564	7.797


  

<b>Ljung-Box (L1) (Q):</b>	0.91	<b>Jarque-Bera (JB):</b>	23.46
<b>Prob(Q):</b>	0.34	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	1.16	<b>Skew:</b>	0.47
<b>Prob(H) (two-sided):</b>	0.56	<b>Kurtosis:</b>	4.45

#### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

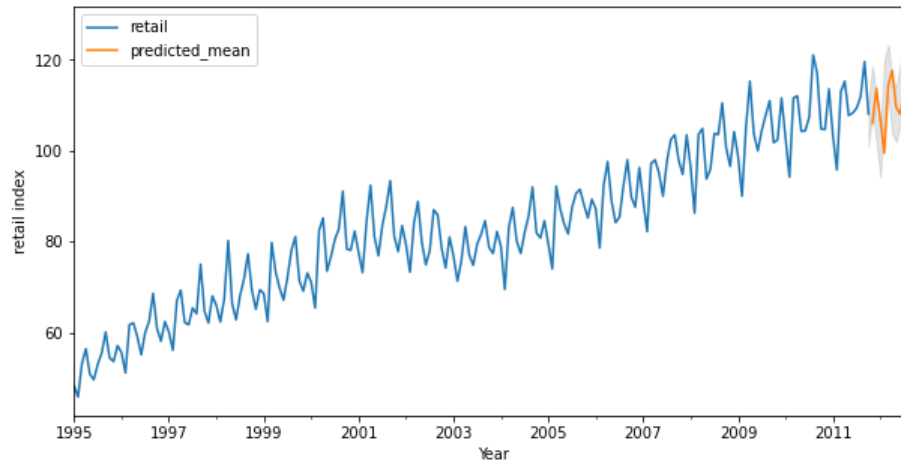
```
In [30]: future_finance_forecast, next_periods_ci = forecast_with_sarimax(retail, 24, order=(0, 1, 1),seasonal_order=(0, 1, 1,12), ylabel = "retail index")
```



by the predicted line we can see that the model is a very good fit to the data



## B

```
In [31]: future_finance_forecast, next_periods_ci = forecast_with_sarimax(retail, 12, order=(0, 1, 10), seasonal_order=(0, 1, 1, 12),  
                                     forecast_next = True, ylabel = "retail index")
```



**C**

```
In [32]: preds, lower_ci, upper_ci = simulate_ets_predictions(retail, "retail",  
                                                           ylabel = 'retail index', xlabel="year", forecast_len  
                                                           = 12)
```



the ES does not work well here because the trend changes, prophet will probably be better.

```
In [33]: retail_prophet = retail.reset_index().rename({"month": "ds", "retail": "y"}, axis =1)
m = Prophet()
m.fit(retail_prophet)
future_dates = m.make_future_dataframe(periods=60, freq='MS')
forecast = m.predict(future_dates)
forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']]
m.plot_components(forecast);
m.plot(forecast);
```

13:03:38 - cmdstanpy - INFO - Chain [1] start processing

13:03:38 - cmdstanpy - INFO - Chain [1] done processing

